

VIJESH REDDY GOLAMARI

+1 (475) 351-5694 | golamarivijesh@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Generative AI Engineer and LLM Specialist with experience designing enterprise AI assistants, RAG systems, agentic workflows, LLM evaluation frameworks, and scalable ML platforms. Delivered a production AI assistant that reduced ticket-resolution time by 40%, expanded foundation-model benchmarking coverage by 45%, and built distributed transformer systems across financial services, technology, e-commerce, and agriculture. Skilled in LLMops, AI safety, model evaluation, multimodal AI, cloud deployment, and cross-functional technical leadership.

TECHNICAL SKILLS

Generative AI & Agentic AI: LLM engineering, prompt engineering, ReAct, function/tool calling, autonomous agents, multi-agent orchestration, memory management, RAG, guardrails, hallucination mitigation, RLHF, DPO

Models & Frameworks: LLaMA 3, GPT-4, Claude, Mistral, BERT, CLIP, LangChain, LangGraph, Hugging Face Transformers, OpenAI APIs, PyTorch, TensorFlow, Scikit-learn

LLMops, MLOps & Evaluation: LLM evaluation and benchmarking, adversarial testing, bias detection, SHAP, distributed training, FSDP, FlashAttention, MLX, Weights & Biases, model monitoring, CI/CD

Data, Search & Platforms: Python, SQL, R, Apache Spark, Kafka, Airflow, Hive, Elasticsearch, BM25, Pinecone, Milvus, PostgreSQL, MySQL, Snowflake, BigQuery, Dataflow

Cloud & Engineering: AWS SageMaker, Lambda, Redshift, Glue, Athena, GCP Vertex AI, Docker, Kubernetes, GitHub Actions, Jenkins, Java, Spring Boot, REST APIs, ReactJS, Django REST Framework, Redis, Jira, Confluence

PROFESSIONAL EXPERIENCE

Symplore Inc

May 2026 - Present

Generative AI Engineer | Project Coordinator | Technical Lead

Michigan

- Lead a cross-functional team of eight engineers across AI, Java backend, frontend, cloud, and DevOps, delivering three AI-enabled features on schedule through Agile planning and Jira-based execution.
- Coordinate sprint planning, delivery timelines, architecture decisions, code reviews, and deployment strategies, enabling 95% of technical deliverables to meet release dates and reducing schedule slippage by two weeks per quarter.
- Design Generative AI applications using LLMs, RAG pipelines, AI agents, recommendation engines, intelligent search, analytics dashboards, and enterprise knowledge systems.
- Build Java and Spring Boot microservices with REST APIs, Kafka, Redis, PostgreSQL, and AWS, supporting 10,000+ concurrent users and reducing API latency by 30%.
- Implement Apache Spark and Dataflow pipelines that improved data-processing throughput by 40%, while mentoring developers across parallel projects.

Agrivision AI

Mar 2026 - Present

Founder & AI Architect

Michigan

- Built and launched YieldAI Global, an AI-powered agriculture intelligence platform providing predictive analytics, crop recommendations, weather intelligence, disease detection, and multilingual digital assistance.
- Owned product development from concept through launch and designed the architecture across ReactJS, Django REST Framework, Python, MySQL, APIs, and cloud infrastructure.
- Developed crop advisory, yield forecasting, irrigation planning, and conversational AI modules, while scaling the platform to support multiple regions and languages.
- Implemented RAG-style retrieval using embeddings, structured prompts, and domain-specific context; integrated Anthropic Claude and OpenAI APIs for grounded, multilingual farmer guidance.
- Evaluated Claude-generated and OpenAI-generated responses against agriculture-specific prompts to improve accuracy, safety, tone, and practical usefulness.

Wipro

Nov 2025 - Apr 2026

GenAI & LLM Evaluation Specialist

Dallas, TX

- Architected scalable LLM evaluation frameworks with Python, Apache Hive, and MLX, increasing benchmarking coverage by 45% across multiple foundation models.
- Designed adversarial and prompt-engineering evaluations with Hugging Face Transformers, increasing detection of jailbreak, hallucination, and multi-hop reasoning failures by 35%.
- Built automated Python and SQL evaluation pipelines that reduced manual validation effort by 60% and accelerated model-release cycles.
- Led AI safety and red-team testing using bias detection and adversarial prompting, reducing harmful outputs by 28% in collaboration with alignment teams.
- Developed ELK-based monitoring dashboards for latency and hallucination trends, reducing model-degradation incidents by 25%, and integrated evaluation into Git and Jenkins CI/CD workflows.
- Applied human-in-the-loop annotation and R-based analysis to assess grounding, instruction following, safety, retrieval quality, reasoning accuracy, latency, and response consistency.

Meta <i>AI/ML Engineer</i>	Aug 2024 - Oct 2025 <i>Menlo Park, CA</i>
• Fine-tuned LLaMA 3 models with PyTorch FSDP on Meta RSC, increasing harmful-content detection by 25% and reducing false positives by 15%. • Developed diffusion and transformer-based generative models that increased ad engagement by 12% and click-through rate by 8%. • Optimized distributed training with FlashAttention and GPU scheduling, reducing training cost by 30% and improving experimentation throughput. • Implemented SHAP and Fairness Indicators pipelines that reduced demographic performance gaps by 22%. • Engineered Spark and Kafka pipelines processing more than 5 TB of daily interaction data and deployed high-throughput models with Docker and Kubernetes.	

Citigroup <i>Applied AI Engineer</i>	Sep 2023 - Aug 2024 <i>San Francisco, CA</i>
• Delivered a Generative AI assistant using GCP Vertex AI and LangChain, reducing HR and IT ticket-resolution time by 40%. • Built RAG pipelines with embeddings, transformer models, and Elasticsearch, reducing hallucinated responses by 30%. • Engineered agentic workflows with API integrations for autonomous task execution across enterprise applications and optimized hybrid BM25/vector retrieval, improving response accuracy by 25%. • Implemented GitHub Actions-based LLMOps pipelines with monitoring and observability, reducing release errors by 35%. • Developed Python and R predictive models that improved credit-risk accuracy by 18%, using statistical analysis and SHAP to identify degradation and guide retraining.	

Flipkart <i>Machine Learning Engineer</i>	Aug 2021 - Nov 2022 <i>India</i>
• Built NLP and embedding models with TensorFlow and Scikit-learn, increasing catalog-prediction accuracy by 15-20%. • Designed recommendation systems using collaborative filtering and feature engineering, driving a 12% uplift in personalized interactions. • Built scalable ML and data pipelines with AWS SageMaker, Redshift, Athena, Spark, Airflow, and AWS Glue, reducing data-preparation time and accelerating large-scale training. • Developed transformer-based conversational systems and optimized inference with batching to improve response accuracy and latency for high-traffic applications. • Implemented Jenkins and GitHub Actions MLOps pipelines, reducing deployment timelines by 40%, and partnered with product and engineering teams to productionize API-integrated ML systems.	

PROJECTS

Autonomous Multi-Agent Enterprise Orchestrator

- Built a LangChain and LangGraph orchestration system for task planning and execution across complex enterprise workflows.
- Implemented Python and SQL tool-calling workflows, stateful agent memory, API integrations, and Airflow orchestration for long-running and parallel processes.

Distributed Multimodal Intelligence Platform

- Developed a PyTorch and CLIP pipeline for joint image-text processing and downstream object-recognition use cases.
- Applied FSDP and FlashAttention for distributed training and engineered Spark pipelines for high-volume ingestion and feature generation.

Enterprise RAG & LLMOps Platform

- Built a LangChain, Elasticsearch, embedding, Hugging Face, and OpenAI API-based retrieval platform for context-aware enterprise document access.
- Implemented GitHub Actions and Weights & Biases workflows for CI/CD, monitoring, deployment stability, and performance tracking.

AI-Powered Jira QA Copilot

- Developed a Python and OpenAI API tool that converts Jira user stories and BRD inputs into structured, automation-ready test cases.
- Integrated Jira REST APIs, NLP parsing, Streamlit, Selenium Java, and Maven to support interactive validation and scalable test automation.

EDUCATION

Sacred Heart University <i>Master's Degree</i>	Jan 2023 - Mar 2024 <i>Fairfield, CT</i>
Loyola Academy <i>Bachelor of Commerce, Business Studies</i>	May 2019 - May 2022 <i>Hyderabad, India</i>

CERTIFICATIONS

- Generative AI with Large Language Models - Coursera
- Generative AI for Data Scientists Specialization - Coursera
- Deep Learning Specialization - Coursera
- Google AI Essentials - Coursera